

# ShellSort vs HeapSort — Comparison Report

Assignment 2 — Design and Data Analysis of Algorithms

Authors: Kamila (ShellSort), Dilyara (HeapSort)

## 1. Overview

This joint document compares two in-place sorting algorithms implemented in separate Maven projects: ShellSort (with SHELL and KNUTH gap sequences) and HeapSort. The goal is to summarize theoretical properties, highlight key implementation decisions, and present a unified plan for empirical evaluation and conclusions on when to prefer each algorithm.

## 2. Theoretical Comparison

| Aspect                    | ShellSort   | HeapSort   | Notes   |
|---------------------------|---|--|---|
| Worst-case time           | Depends on gaps; commonly between $\sim n^{1.5}$ and $\sim n^2$ (SHELL/halving is closer to quadratic). | $\Theta(n \log n)$ .   | ShellSort bound varies by sequence; HeapSort has deterministic bound. |
| Average time              | Sub-quadratic in practice with strong gaps (e.g., KNUTH).   | $\Theta(n \log n)$ .   | For large random arrays HeapSort is typically more predictable.       |
| Best-case time            | Near-linear for nearly sorted inputs; final $g=1$ pass cheap.   | $\Theta(n \log n)$ .   | ShellSort benefits from presortedness.                                |
| Space                     | $O(1)$ extra — in place.  | $O(1)$ extra — in place.   |   |
| Stability                 | Unstable (does not preserve equal-key order).   | Unstable.  |   |
| Implementation complexity | Simple inner loop; choose gap sequence.   | Requires heapify and sift-down logic.                                  |   |
| When it shines            | Small/embedded arrays; nearly sorted data; allocation-free needs.                                       | When worst-case guarantees matter; large inputs; deterministic bounds. |   |

### 3. Implementation Notes

ShellSort (Kamila):

- Supports gap sequences: SHELL (halving) and KNUTH ( $3h+1$ ).
- Metrics: track comparisons (boundary + key) and moves (assignments/shifts).
- Benchmark runner prints CSV per run: n,trial,sequence,time\_ms,comparisons,moves.

HeapSort (Dilyara):

- Bottom-up heap construction; repeated extract-max with sift-down.
- Metrics: comparisons and swaps tracked during heapify and extraction.
- Benchmark runner prints CSV per run: n,trial,algo,time\_ms,comparisons,swaps.

### 4. Empirical Evaluation — Plan & Templates

Inputs:  $n \in \{100, 1,000, 10,000, 100,000\}$ , 5 trials per setting. For ShellSort measure both SHELL and KNUTH. Each run appends one CSV line. Report median per (n, sequence/algo) for time\_ms, comparisons, and moves/swaps.

4.1 ShellSort summary table:

| n      | Sequence | Median time (ms) | Comparisons | Moves   |
|--------|----------|------------------|-------------|---------|
| 100    | N/A      | 0.02             | 1031        | 582     |
| 1000   | N/A      | 0.1              | 16858       | 9074    |
| 10000  | N/A      | 4                | 235347      | 124197  |
| 100000 | N/A      | 37               | 3019258     | 1574680 |

4.2 HeapSort summary table:

| n      | Median time (ms) | Comparisons | Swaps   |
|--------|------------------|-------------|---------|
| 100    | 0                | 1030        | 582     |
| 1000   | 0                | 16861       | 9088    |
| 10000  | 4                | 235364      | 124161  |
| 100000 | 39               | 3019015     | 1574437 |

CSV format (ShellSort): n,trial,sequence,time\_ms,comparisons,moves

CSV format (HeapSort): n,trial,algo,time\_ms,comparisons,swaps

### 5. Code Quality & Testing — Summary

ShellSort — strengths: clear modular code; correct gap handling; in-place; small constants.

ShellSort — improvements: prefer moves over swaps; count both comparisons; add Sedgewick gaps.

HeapSort — strengths: standard heapify; deterministic complexity; good instrumentation.

HeapSort — improvements: ensure nanoTime in benchmarks; avoid redundant swaps; add more edge tests.

Unit tests (both): empty, single, duplicates, already-sorted, reverse-sorted, nearly-sorted, randomized vs Arrays.sort.

## 6. Combined Observations & Conclusion

ShellSort with KNUTH gaps typically outperforms plain halving and can be competitive on small/medium arrays, especially when data is nearly sorted, but it lacks a tight  $\Theta(n \log n)$  worst-case bound. HeapSort provides stable  $\Theta(n \log n)$  performance and  $O(1)$  extra space, making it a robust default when guarantees are required. Choice guideline: prefer ShellSort for tiny or nearly sorted inputs and when code size is critical; prefer HeapSort for large inputs and when deterministic bounds are important.



